SCORE: A Framework for Synthesizing Complete, Owned, Responsive Environments

J.P. Kardolus, Babon Innovations B.V., Utrecht, The Netherlands

Working draft, 13 October 2026

Abstract

Games, films and robot training need 3D worlds that are complete, editable and owned by their makers; world models and one-shot generators give frames no engine opens, or one fused scene without separate objects or clear rights. Studios build worlds in stages, from concept art to review. SCORE keeps these stages and automates them with open vision foundation models, procedural tools and coding agents, while the creator chooses and reviews. Every stage is a node of one runner: its output is keyed on its inputs and code, judged by gates that answer pass, fail or unknown, and accepted only when nothing got worse, so a fault is fixed in its own stage and only what follows it runs again; paid steps are cleared against a budget read from the provider's bill. Every model allows commercial use of its output, so the creator stays the author. We describe the framework as built for its first world, the game world of 2099 with 18 places, and the measured gates of its surfaces, resting, light, motion, effects, sound and people. Code and paper are available at https://github.com/Babon-Innovations-b-v/score (MIT license).

Gap: Figure 1, four worlds. World 1's square, hub and Mars camp beside their concepts; worlds 2 to 4 as not started. It waits on those places' completion gates.

1 Introduction

Interactive 3D worlds are the raw material of games, film and embodied AI [1], and building one is expensive: every object is modelled, surfaced, placed, given collision, lit and given sound in one style, and the world stays editable. Generative models promise to take over much of this, but in GDC's 2026 survey 52 percent of game workers thought generative AI was hurting the industry [2]; a tool for building worlds has to leave the result in the creator's hands.

World models generate a world as frames [3, 4] that no engine can open. 3D world generators [5, 6] produce geometry, but mostly as one fused asset, from which separate objects are recovered after the fact [7, 8]. Both work in one shot, so a creator cannot find or change the step behind a fault. Procedural systems [9, 10] and coding agents that write scenes as programs [11–14] keep the world explicit, but rebuild a scene that already exists or draw their objects from a fixed library, and such agents judge their own geometry unreliably [15]. Open vision foundation models [16] now read shapes [17], masks [18], depth [19] and parts [20] out of pictures, and a coding agent can connect them as tools; none of them provides one style across objects that different models made, which the field names as an open problem [21, 22].

Studios build worlds in stages, each with an output a person can check; SCORE keeps these stages and automates them. Its worlds are **complete**: every object, surface, light and sound exists wherever the player looks. They are **owned**: every model in the framework allows commercial use of its output, so the creator stays the author, holding the full commercial rights. They are **responsive**:

every object has real collision, mass and material. And they can be any kind of **environment**. To our knowledge, no other openly available world generation work joins vision foundation models into engine-ready scenes that are built from reference pictures, steered by a person's choices at each step, and made only with models that allow commercial use of their output.

We contribute **the SCORE framework**, stages as nodes of one runner that keys, judges and accepts every output over one canonical scene (Section 2); **measured gates** for every layer of a place and a completion gate over them (Section 3); and **evidence from a first world**, with its spend per stage (Section 4).

2 The framework

A world is a table of stages; world 1's has 95 rows. A **node** is one stage for one place (`assembly@hangar`) or for the world, and one runner makes, judges and accepts every node. The creator's picks (concepts, style, references) are the inputs; from there everything runs without a typed yes, and the creator reviews.

2.1 Keys, gates and accepts

Keys. A node is made only when its key has no finish record. The key hashes the stage and place, a code version, the contracts of its output and inputs, its parameters, a seed derived from the world's, the pinned model images, every world file it reads and the accepted output of every node it reads. The code version hashes every source file the stage imports, leaving out the runner and the renting code, which change what a run costs, not what it makes. A fault fixed in one stage moves only its key and those after it. Outputs are kept by content hash, never overwritten.

Gates and accepts. Each gate answers pass, fail or unknown, with measures and a reason; missing evidence is unknown, and unknown blocks like fail. The runner accepts an output by itself when every gate passed and it is no worse than the one accepted before: no answer dropped, no measure moved the wrong way by a step a viewer could see, no gate missing. Later nodes read only accepted outputs. The creator keeps four commands, from his own terminal: reject an accepted output (it moves back to the one before, and every node that read it is stale), accept a worse one with his reason, waive one gate for one key, and override the runner for one paid run. None is needed for a run; a barrier keeps agents from them (a terminal, a store that refuses the creator's records from an agent's process, and a hook).

Spending and review. A paid node is cleared only when every gate on its inputs passed and the month's ceiling has room for its estimate (Section 3.7); everything that rents a machine or calls a paid model refuses without that ticket, and buying again more than EUR 5 of accepted outputs needs a written reason. Each node's review page is built from the store alone, and the creator reviews a place in a USD viewer on the store's copy of its accepted assembly; nothing is rendered for review.

2.2 From concept to scene

Inventory, plan and record. A coding agent lists every visible element of the concept as an inventory row (kind, anchor, size, count) or drops it with a written reason; nothing is made or placed that is not a row. It draws a dimensioned plan, gated on `mapped` (every element a row, a door or a written drop; every door says where it leads). The scene record is written from the plan, inventory, made ground and sky, gated on `game-free` (nothing copied from the game the world

came from), `plan-matched` (the shell within 5 cm of the plan), `openings` (every doorway onto a floor, no view into nothing) and `seed-matched` (every old entry matched or dropped with a reason).

The canonical scene. Every place is one OpenUSD stage [23]; each object, one per real-world object, carries its kind, transform, collision and each part's library surface. Its layers lie in a fixed order, the creator's edit layer on top, then people, sound, light, motion, effects, crops, machines, settle, ground and the generated base, so an edit survives a regenerated base. Each place is assembled once, from every layer's accepted output; its `layers` gate answers unknown while a needed layer has none, and an old output never stands in.

Models, generation first. Every object goes the prop route first: Nano Banana Pro [24] draws its close-up from its crop of the concept, Pixal3D [17] builds it, the surface stage paints it, and a likeness gate compares its outline from eight sides with its close-up. Only an object whose checks fail is built by a code builder a coding agent writes, reversing the earlier parts check (the creator's choice, 12 October 2026). On 13 of the hub's objects both orders end the same (Table 1), but generation first draws a close-up of every object. Every close-up of eight thin and flat objects failed the shape check, so code builds carry them. The hatch wheel shows the gain: its code build's outline overlap was 0.45, a fail, and its generated model's 0.60. Most of world 1's models predate the change and enter as unverified seeds.

Table 1. How 13 hub objects, chosen to cover the kinds, end in each order. Close-ups cost EUR 8.12 and the sample's prop route EUR 13.80, about EUR 2.20 an object; no take was baked (the run stopped at the keyboard's paint gate).

Order	Prop route	Code build	Needs a new builder
Generation first	4	7	2
Code first	4	7	2

Gap: Figures 2 to 5. (2) The stages on the hub, from concept to assembled room; (3) the greenhouse's review page; (4) the wreck with an edit that survives a regenerated base whose models were repainted; (5) four of Table 1's objects from close-up to final model. They wait on those places' assemblies and on renders of the sample's takes; the earlier hub figure predates the runner and is withdrawn, and the comparison recorded no agent hours.

3 Stages and their gates

3.1 Surfaces: the finish map

Models give shape only; every part takes one surface from a shared library of rule-based surfaces, written as ProcFunc functions [25] on Infinigen's shaders [9], coloured from the place's palette and worn only where wear has a cause (Appendix D). Gemini 3.1 Pro [26] lists the close-up's finishes and names a library surface for each. Nano Banana Pro draws the close-up again with one flat colour per finish, up to three times until the map fits the close-up's outline, then draws unlit renders of the raw model's sides and back against it. SegviGen [27] splits the mesh guided by the close-up's map, every map votes on the faces its view sees, and a part whose second finish holds a fifth of it fails. Each part is painted from its library surface and baked, never in the picture's colours, which gave every piece its own rust and stains (Appendix A).

The bake is checked against the close-up. A finish's surface must lie within ΔE 32 of its colour there, else a nearer surface of its family is taken or a new one is proposed to the creator, and the part fails by name. From the greenhouse's water tank, which the creator found baked as three near greys and tipped, came two checks: a bake fails when it darkens a finish the close-up shows lighter,

or turns all surfaces into one grey; and short legs are stretched onto one level plane, a foot over 3 cm off the floor failing. Dirt follows the close-up.

3.2 Resting

Loose objects are dropped onto the made ground in a physics simulation; one that would turn more than 15 degrees or drift more than 30 cm is a resting fault of the layout and keeps its laid pose. The settle stage lays where each object comes to rest as a layer (held objects snapped onto their holders, a stack onto one support, an object that would rest worse kept as laid), gated on `rested` and `resting` (no fault that appears with the layer and not without it).

3.3 Light

Lamp strengths are measured, never copied from the source game. The light layer places live lamps from the scene record, with the sky, the sun and the Earth. One paid render per place gives every lamp its own Cycles light group on the place's views; because light adds up, the look at any strengths is mixed from these passes without rendering again, and a free sweep chooses the scale that brings the place's brightness within 1.5 times its concept's, halving any lamp that turns what it lights white. Outdoors the suns are swept after the lamps. A lamp whose fixture stops 90 percent of its rays moves out of it, and one behind a clear lens stands just outside, as Cycles passes no light through a glowing lens.

3.4 Motion, effects and sound

Motion. Each mover (a path, a hinge, a spin, a rail, a turn, a climb) has its kind's speed range cited from real machines, gated on `clear` (no mover within 3 cm of a still thing or another mover, read every quarter second), `speeds`, `loops`, `grounded` and `covered` (every still prim of a moving kind says why). **Effects** are emitters measured against real footage (photographs of lift-offs, rover films of Martian dust devils), gated on `attached`, `footage`, `light` (the flame lights the pad at least twice as brightly as the place's lamps) and `outside` (no open-sky particle inside a room). **Sound:** every source a place names (room tones, objects, motions, each surface's steps) gets the world's picked file or the best of eight MOSS-SoundEffect [28] takes of its brief, gated on `sources`, `level`, `likeness` (CLAP [29] score at least 0.30, the tenth percentile of the world's picks) and `no freesound`.

3.5 People

People are made like props: a FLUX.2 klein [30] close-up in SOMA-X's bind pose, built by Pixal3D, painted by the surface line, over SOMA-X's body fitted just inside (Gauss-Newton, 24 shape numbers), whose skin weights each look point takes from its closest triangle; Kimodo's clips [31], from sentences, move it. The **shell gate** passes a look when 99 percent of its soft points lie within 5 cm of the body, its worst point recorded. Hard parts ride one joint, belts are cloth, each clip gets its own verdict (tucks, hard parts stretched, cloth or body through cloth) and casts use only passing clips; a far-only look is judged on its far body (Appendix B). **Casts:** SAM 3 [18] finds the concept's figures, a judge tells people from pictures, reflections and empty suits, and casts are gated on `counted`, `standing`, `clear`, `made` and `drapes`; no crowd copy fills a tenth of a view's height.

3.6 Completion

A place is complete when its gate passes on its final assembly: every row made or dropped with a reason, nothing unmade, no placeholder or resting fault, the creator's picks shown, every layer

and shown place present, nothing from the source game. A hook stops an agent claiming otherwise: agents judging their own scenes agree with a measured geometry score only 45.8 percent of the time [15].

Gap: Figures 6 to 10. (6) The water tank painted, the tool board in picture against library colours; (7) the lit Earth, a swept room; (8) lift-off, a rover, a dust devil, an airlock; (9) the player walking, the crowd; (10) a room in a physics engine. They wait on assemblies; until the habitat's physics trial passes, responsiveness is that trial.

3.7 Budget and ledgers

Machines and Google calls write ledger lines with their node's ticket; the month's spend is the higher of ledgers and bill, which caught misses: EUR 1,554 billed against EUR 1,387 in the cloud ledger on 12 October 2026, and EUR 549 billed by Google for a day the earlier ledger put at EUR 30. Over 13.8 hours of world 1's run the ledgers hold EUR 117.45 of machines and EUR 34.19 of Google calls, most for painting (EUR 37.61), the finish maps (EUR 26.59), the people's bodies (EUR 16.09) and the light passes (EUR 10.17).

4 Evidence

Four worlds are planned; world 1's 18 places are judged once their layers stand accepted.

Table 2. People, 13 October (experimental, Section 3.5): usable clips of 19 per look by the gate as built, on drawn surfaces only, and as one piece on what cameras see. Faults and fixes: Appendix B.

Look	Body gates	Gate	Visible	Cameras, one piece
Player (work)	pass	5	0	0
Bram (work)	pass	5	0	not run
Sefa (work)	pass	0	0	not run
Nev (work)	pass	0	0	not run
Kit, average man (olive)	fail: triangles	0	not run	not run
Player, wider pose (work)	fail: triangles	0	not run	0
Kit, average man, wider pose (olive)	fail: shell	0	not run	0

Gap: Table 3 and the film. Per place: rows, models made and coded, layers, faults before and after spending, time, euros, verdict; a walk through world 1. They wait on the completion gates.

Agent tools. With them an agent used about half the output and cost on two of three places; on the third it left more faults; one run each (Appendix E).

5 Limitations, future work and conclusion

One creator judges; he accepted the hub before the runner. Closed models (Nano Banana Pro, Gemini 3.1 Pro) may change or go. Hard belts failed deep bends; light is live, measured on earlier layouts; objects are not interactive. Next: three worlds, an engine adapter, interaction, baked lamps, open models, locked parts, LEGO-Bench [15]. In short, SCORE keeps, keys and judges every stage's output, so a fault is fixed where it happens; completed places will show if it holds.

References

[1] Yue Yang, Fan-Yun Sun, Luca Weihs, Eli VanderBilt, Alvaro Herrasti, Winson Han, Jiajun Wu, Nick Haber, et al. Holodeck: Language guided generation of 3D embodied AI environments. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2024. arXiv:2312.09067.

[2] Diego Argüello. 36 percent of game workers use generative AI, but half think it's bad for the industry. Game Developer, February 2026. Reporting GDC's 2026 State of the Game Industry survey.

[3] Jake Bruce, Michael Dennis, Ashley Edwards, Jack Parker-Holder, et al. Genie: Generative interactive environments. *arXiv preprint arXiv:2402.15391*, 2024.

[4] Jack Parker-Holder and Shlomi Fruchter. Genie 3: A new frontier for world models. Google DeepMind blog, August 2025. URL https://deepmind.google/blog/genie-3-a-new-frontier-for-world-models/. Vendor material.

[5] Dilin Wang, Hyunyoung Jung, Tom Monnier, Kihyuk Sohn, et al. WorldGen: From text to traversable and interactive 3D worlds. *arXiv preprint arXiv:2511.16825*, 2025.

[6] Tencent Hunyuan. HY-World 2.0. GitHub repository, 2026. URL https://github.com/Tencent-Hunyuan/HY-World-2.0. Tencent HY-World 2.0 Community License; not applicable in the EU, UK and South Korea.

[7] Jichen Hu, Jiawei Guo, Jiazhong Cen, Chen Yang, Sikuang Li, and Wei Shen. WorldAct: Activating monolithic 3D worlds into interactive-ready object-centric scenes. *arXiv preprint arXiv:2605.15843*, 2026.

[8] Muyao Niu, Jixuan He, Ruihan Yu, Lian Fu, et al. WorldSculpt: Generating compositional worlds from grounded videos. *arXiv preprint arXiv:2609.05416*, 2026.

[9] Alexander Raistrick, Lahav Lipson, Zeyu Ma, Lingjie Mei, Mingzhe Wang, Yiming Zuo, Karhan Kayan, Hongyu Wen, Beining Han, Yihan Wang, Alejandro Newell, Hei Law, Ankit Goyal, Kaiyu Yang, and Jia Deng. Infinite photorealistic worlds using procedural generation. *arXiv preprint arXiv:2306.09310*, 2023.

[10] Alexander Raistrick, Lingjie Mei, Karhan Kayan, David Yan, Yiming Zuo, Beining Han, Hongyu Wen, Meenal Parakh, Stamatis Alexandropoulos, Lahav Lipson, Zeyu Ma, and Jia Deng. Infinigen Indoors: Photorealistic indoor scenes using procedural generation. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2024. arXiv:2406.11824.

[11] Ziniu Hu, Ahmet Iscen, Aashi Jain, Thomas Kipf, et al. SceneCraft: An LLM agent for synthesizing 3D scene as Blender code. In *International Conference on Machine Learning (ICML)*, 2024. arXiv:2403.01248.

[12] Zhiqi Li, Yuxuan Liao, and Bo Zhu. Recursive code world models: Building complex worlds through recursive scene programs. *arXiv preprint arXiv:2609.11499*, 2026.

[13] Chunchao Guo, Jinpeng Li, Yang Li, and Zilong Huang. WorldClaw: Agentic 3D open-world generation at scale. *arXiv preprint arXiv:2608.05248*, 2026.

[14] Lei Yin, Wentao Cheng, Zhida Qin, Tianyu Huang, et al. AutoUE: Automated generation of 3D games in Unreal Engine via multi-agent systems. *arXiv preprint arXiv:2603.07106*, 2026.

[15] Xirui Li, Peng Shi, Mingwen Dong, Sheng Zhang, Zhuoyan Xu, Dongkyu Lee, Shuaichen Chang, Yi Xiang, Lin Pan, and Jiarong Jiang. LEGO-Anything: Coding agents for 3D scene reconstruction. *arXiv preprint arXiv:2609.36380*, 2026.

[16] Maxime Oquab, Timothée Darcet, Théo Moutakanni, et al. DINOv2: Learning robust visual features without supervision. *arXiv preprint arXiv:2304.07193*, 2023.

[17] Dong-Yang Li, Wang Zhao, Yuxin Chen, Wenbo Hu, Meng-Hao Guo, Fang-Lue Zhang, Ying Shan, and Shi-Min Hu. Pixal3D: Pixel-aligned 3D generation from images. *arXiv preprint arXiv:2605.10922*, 2026.

[18] Nicolas Carion et al. SAM 3: Segment Anything with concepts. *arXiv preprint arXiv:2511.16719*, 2025.

[19] Ruicheng Wang et al. MoGe-2: Accurate monocular geometry with metric scale and sharp details. *arXiv preprint arXiv:2507.02546*, 2025.

[20] Yuchen Lin, Chenguo Lin, Panwang Pan, Honglei Yan, Yiqiang Feng, Yadong Mu, and Katerina Fragkiadaki. PartCrafter: Structured 3D mesh generation via compositional latent diffusion transformers. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2025. arXiv:2506.05573.

[21] Jiafeng Wu, Zhuofan Lou, Jian Liu, Dazhao Du, Chunchao Guo, and Song Guo. From visual synthesis to interactive worlds: Toward production-ready 3D asset generation. *arXiv preprint arXiv:2604.23629*, 2026.

[22] Zhifei Yang, Guangyao Zhai, Keyang Lu, YuYang Yin, Chao Zhang, Zhen Xiao, Jieyi Long, and Nassir Navab. FlowScene: Style-consistent indoor scene generation with multimodal graph rectified flow. *arXiv preprint arXiv:2603.19598*, 2026.

[23] Pixar Animation Studios and Alliance for OpenUSD. OpenUSD: Universal scene description, 2026. URL https://openusd.org/.

[24] Google. Gemini 3 Pro Image ("Nano Banana Pro") through the Gemini API, 2026. URL https://ai.google.dev/. Paid service. Vendor material.

[25] Alexander Raistrick, Karhan Kayan, Jack Nugent, David Yan, Lingjie Mei, Meenal Parakh, Hongyu Wen, Dylan Li, Yiming Zuo, Erich Liang, and Jia Deng. ProcFunc: Function-oriented abstractions for procedural 3D generation in Python. *arXiv preprint arXiv:2604.26943*, 2026.

[26] Google. Gemini 3.1 Pro (gemini-3.1-pro-preview) through the gemini API, 2026. URL https://ai.google.dev/. Paid service. Vendor material.

[27] Lin Li, Haoran Feng, Zehuan Huang, Haohua Chen, Wenbo Nie, Shaohua Hou, Keqing Fan, Pan Hu, Sheng Wang, Buyu Li, and Lu Sheng. SegviGen: Repurposing 3D generative model for part segmentation. *arXiv preprint arXiv:2603.16869*, 2026.

[28] OpenMOSS-Team. MOSS-SoundEffect v2.0. Hugging Face, 2026. URL https://huggingface.co/OpenMOSS-Team/MOSS-SoundEffect-v2.0. Apache-2.0.

[29] Yusong Wu, Ke Chen, Tianyu Zhang, et al. Large-scale contrastive language-audio pretraining with feature fusion and keyword-to-caption augmentation. *arXiv preprint arXiv:2211.06687*, 2022.

[30] Black Forest Labs. FLUX.2 [klein] 4B model card. Hugging Face, 2026. URL https://huggingface.co/black-forest-labs/FLUX.2-klein-4B. Apache-2.0. Vendor material.

[31] Davis Rempe, Mathis Petrovich, Ye Yuan, Haotian Zhang, Xue Bin Peng, et al. Kimodo: Scaling controllable human motion generation. *arXiv preprint arXiv:2603.15546*, 2026.

[32] Maureen Stone, Danielle Albers Szafir, and Vidya Setlur. An Engineering Model for Color Difference as a Function of Size. In *Color and Imaging Conference (CIC 22)*. Society for Imaging Science and Technology, 2014. URL https://library.imaging.org/admin/apis/public/api/ist/website/downloadArticle/cic/22/1/art00045.

[33] Blender Foundation. Blender, 2026. URL https://www.blender.org/. Version 5.0.1 used here. GPL.

[34] Qwen Team. Qwen-Image-Edit-2511 model card. Hugging Face, 2025. URL https://huggingface.co/Qwen/Qwen-Image-Edit-2511. Apache-2.0. Vendor material.

[35] Qwen Team. Qwen3.8-27B model card. Hugging Face, 2026. URL https://huggingface.co/Qwen/Qwen3.8-27B. Apache-2.0. Vendor material.

[36] Oriane Siméoni, Huy V. Vo, Maximilian Seitzer, et al. DINOv3. *arXiv preprint arXiv:2508.10104*, 2025.

[37] Ken Deng, Yunhan Yang, Jingxiang Sun, Xihui Liu, Yebin Liu, Ding Liang, and Yan-Pei Cao. GeoSAM2: Unleashing the power of SAM2 for 3D part segmentation. *arXiv preprint arXiv:2508.14036*, 2025.

[38] Lihe Yang, Bingyi Kang, Zilong Huang, et al. Depth Anything V2. *arXiv preprint arXiv:2406.09414*, 2024.

[39] World Labs. Terms of Service, 2026. URL https://docs.worldlabs.ai/terms-of-service. Last updated 21 January 2026; accessed 2026-10-08. Vendor material.

[40] Parishad BehnamGhader, Vaibhav Adlakha, Marius Mosbach, Dzmitry Bahdanau, Nicolas Chapados, and Siva Reddy. LLM2Vec: Large language models are secretly powerful text encoders. In *Conference on Language Modeling (COLM)*, 2024. arXiv:2404.05961.

[41] Aaron Grattafiori, Abhimanyu Dubey, Abhinav Jauhri, et al. The Llama 3 herd of models. *arXiv preprint arXiv:2407.21783*, 2024.

[42] Xitong Yang, Devansh Kukreja, Don Pinkus, others, and Kris Kitani. SAM 3D Body: Robust full-body human mesh recovery. *arXiv preprint arXiv:2602.15989*, 2026.

[43] Ammar Qammaz. SAM3DBody-cpp: a standalone C++ inference engine for SAM 3D Body. GitHub, 2026. URL https://github.com/AmmarkoV/SAM3DBody-cpp. MIT code; its converted weights are under the SAM Licence.

[44] Maria Korosteleva and Olga Sorkine-Hornung. GarmentCode: Programming parametric sewing patterns. *ACM Transactions on Graphics*, 42(6):197, 2023. arXiv:2306.03642.

[45] The Newton Developers. Newton: a GPU-accelerated physics simulation engine, 2026. URL https://github.com/newton-physics/newton. Version 1.6.1 used here (SolverVBD). Apache-2.0.

[46] Miles Macklin. Warp: A high-performance Python framework for GPU simulation and graphics, 2022. URL https://github.com/NVIDIA/warp. Version 1.18.0 used here. Apache-2.0 since 1.6.2.

[47] Chongjie Ye, Yushuang Wu, Ziteng Lu, Jiahao Chang, Xiaoyang Guo, Jiaqing Zhou, Hao Zhao, and Xiaoguang Han. Hi3DGen: High-fidelity 3D geometry generation from images via normal bridging. In *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, 2025. arXiv:2503.22236.

[48] Black Forest Labs. FLUX.2 [klein] Base 4B model card. Hugging Face, 2026. URL https://huggingface.co/black-forest-labs/FLUX.2-klein-base-4B. Apache-2.0. Vendor material.

[49] thedeoxen. refcontrol: reference-depth LoRA for FLUX.2 [klein] 4B. Hugging Face, 2026. URL https://huggingface.co/thedeoxen/refcontrol-FLUX.2-klein-4B-reference-depth-lora. Apache-2.0; training data not disclosed.

[50] Jia-Peng Zhang, Cheng-Feng Pu, Meng-Hao Guo, Yan-Pei Cao, and Shi-Min Hu. One model to rig them all: Diverse skeleton rigging with UniRig. *ACM Transactions on Graphics (SIGGRAPH)*, 2025. arXiv:2504.12451.

[51] Rinat Abdrashitov, Kim Raichstat, Jared Monsen, and David Hill. Robust skin weights transfer via weight inpainting. In *SIGGRAPH Asia 2023 Technical Communications*, 2023.

[52] Anthropic. Claude Code, 2026. URL https://www.anthropic.com/claude-code. Paid subscription. Vendor material.

[53] Hou In Ivan Tam, Hou In Derek Pun, Austin T. Wang, Angel X. Chang, and Manolis Savva. SceneEval: Evaluating semantic coherence in text-conditioned 3D indoor scene synthesis. *arXiv preprint arXiv:2503.14756*, 2025.

[54] Hongchi Xia, Xuan Li, Zhaoshuo Li, Qianli Ma, Jiashu Xu, Ming-Yu Liu, Yin Cui, Tsung-Yi Lin, Wei-Chiu Ma, Shenlong Wang, Shuran Song, and Fangyin Wei. SAGE: Scalable agentic 3D scene generation for embodied AI. *arXiv preprint arXiv:2602.10116*, 2026.

Appendix

A The hub in six rounds

The surface library, one model per real-world object and the straightness check came out of six rounds of work on the hub on 6 and 7 October 2026, each answering the creator's verdict on the round before (Table A1). In the first two rounds, he chose between the old and the new method side by side from the same cameras.

Table A1. The hub's six rounds. Cost is given as rented GPU time in euros and picture-model calls in dollars. Frame time is the mean frame time on an RTX 5080, in the night scene, at 1600×900.

Round	Change	Creator's verdict	Cost	Frame time
1, two walls, A/B	surface library, routing by shape, room checks	chose the new method ("Y looks much cleaner")	€0.88	7.81 ms vs 7.53
2, two walls, A/B	71-variant library, shared texture sets, roof check	chose the new method again	€1.70	8.05 ms vs 8.24
3, whole room	all stages end to end, 74 models	old models were still loaded; labels and floor fittings were wrong	€1.67	6.78 ms (old 9.22)
4, whole room	detail through the pipeline, picture colours kept	"pretty much all of this became messy"	€8.50, $5.09	9.6 ms
5, eight pieces, A/B	shape only from the pipeline, library surfaces, parts check	new method preferred for the door, notice board and porthole; three pieces still wrong	€0.21	not run
6, whole room	round 5 everywhere, one model per object, straightness check	"this looks amazing, I have no negative feedback"	€3.48, $1.07	5.6 ms

Three findings shaped the method. First, deterministic checks catch what agents miss: in round one, the share of rays leaking out through walls and roof fell from 1.53 percent of 5,817 to zero, and the colour noise within surfaces from a median of 31.1 to 3.2 percent, but only walking through the loaded room showed 26 old models still in it, which became the made-only check. Second, generated pieces that kept their pictures' colours made the room inconsistent, so surfaces come only from the library. Third, generated composite objects melted their parts together, and generated boxes leaned (one locker by 13 degrees), which led to one model per real-world object and to the straightness check. By round six the room used 0.45 million triangles, against 1.96 million in round four, and 187 MB on disk, against 487 MB.

B Check details

- **Room checks** cast rays to find light leaking out through walls and roof, find pieces in front of a real opening and roof and wall fittings more than 3 cm off their surface or turned more than 5 degrees from it, and count pieces, triangles and lamps against the room's budget.
- **Door check:** every door must stand in a wall or partition that fully separates its two sides, tested by walking a grid from one side to the other.
- **Model check:** every model must have walls at least 3 mm thick. A generated box fails if one of its sides tilts more than 2 degrees or is warped. Making a generated model solid stops if its bounding box grows more than 3 percent.
- **Likeness gate:** a made model's outline from eight sides must overlap its close-up's by at least 0.50, a least set on the hub's twelve accepted generated models with a margin of 0.05, and its DINOv2 [16] similarity must reach a least set the same way; without that calibration the gate answers unknown.
- **Finish maps:** a close-up's map passes when its outline overlaps the close-up's by at least 0.85 and it uses every listed finish; a part whose second finish holds 20 percent of its voted area mixes finishes and fails.
- **Bake check:** each baked part within ΔE 20 (CIE76) of its surface's colour and within ΔE 32 of its close-up's; lightness order fails when the close-up shows one surface more than 10 L* lighter and the bake makes it more than 3 L* darker; one look fails when no two near-neutral surfaces bake far enough apart for a viewer to tell at the place's view distance, while

the close-up shows them more than ΔE 24 apart. What a viewer tells is measured, not set: Stone, Szafir and Setlur measured the CIELAB difference half of their viewers see between two colours s degrees wide, per axis ND(50, s) = C + K/s, with (C, K) = (5.08, 0.75) for L*, (5.34, 1.54) for a* and (5.35, 2.87) for b* on sizes 1/3 to 6 degrees [32]; two surfaces are told apart when their difference, each axis over its ND, is 1 or more. A surface's size is its model's middle size times the root of its area share, seen from the place's nearest view. This replaced a fixed ΔE 12, which failed the wreck's access cover on a 9.3 L* step between its steel and its scorched white that its owner sees (11.3 degrees from the nearest view; half of viewers see 5.2 L* at that size); a quadrant of a standing model's base more than 3 cm off the floor fails. Glass, screens, lamps and print are judged by their own colour and left out of the order.

- **Settling and resting check:** an object that would turn more than 15 degrees or drift more than 30 cm when dropped keeps its laid pose and is marked as not resting; small debris, with no side over 1.4 m, may turn freely and drift up to 1 m. Hung objects are exempt by their inventory anchor, and fixed objects are not judged for tipping. The triage fails a foot off its support and an up axis leaning more than 3 degrees.
- **Light:** the place's median brightness within 1.5 times its concept's and every view within 2.5 times; a lamp is blown when over 20 percent of the pixels where it gives at least half the light are white, and unlit when it gives less than 0.001 percent of every view's light.
- **Motion:** no mover within 3 cm of a still thing or another mover, read every 0.25 s; rounds join within 1 mm; speeds within each kind's cited range.
- **Sound:** CLAP score against the brief at least 0.30, the tenth percentile of 55 of the world's picks; a take no more than 0.5 s longer than its brief.
- **People:** the shell gate passes a look when 99 percent of its soft points lie within 5 cm of the fitted body; the worst point and its joint are recorded. A clip is usable when no frame tucks more than 5 mm, no hard part stretches past 1.2 times its bind length, cloth crosses cloth no more than in the earlier body's same clip, and no part's edges pass through cloth deeper than 3 mm more often than at the bind. A far-only look is judged on its joins, likeness and the far body's triangles (at most 1,220); the near outfit's triangles and shell offset are recorded, not gating. The A-pose outline is judged against the finished Pixal3D take the build starts from; the take's own loss against the raw model is recorded apart. A crowd copy stands only where it fills under a tenth of every view's height.
- **Worse than accepted:** a measure counts as worse only when it moves the wrong way by at least its tolerance, the smallest change a viewer could see: one more fault for a count; 0.02 of full scale for lightness on 0 to 1 (a halo, fill, beam or haze wash), and 5 of 255 levels for a colour gap, the step a viewer notices between two pictures; 0.05 for a brightness ratio; one percent of the picture for a share of it; a twentieth of the Earth's lit disc; 5 mm for a length. A score no viewer reads directly, such as a likeness score, is worse by any move. Tolerances decide only whether an output waits for the creator; a check that fails still fails.
- **The creator's accept carried forward:** when he accepts a worse output, his accept records what he allowed, each measure's value before and after. A later output that passes every gate and is no worse than the output he took, within the tolerances, is accepted by the runner with his accept named; a move beyond it still waits for him.
- **What the creator acts on:** he rejects, accepts or waives only an output he has seen on its review page. This is his rule; the commands do not yet check that the page was opened.
- **People seen by the cameras (experimental):** a clip's crossings of body through cloth are counted only where a camera that will show the person can see them: the place's views at the person's cast spot, and a close front and three-quarter view. Crossings no camera sees are recorded, not counted against the clip. So far this is a measure on the stored bodies (Section

3.5); the per-clip gate does not use it yet.

- **Completion gate:** the requirements given in Section 3.6, each with its result tied to a fingerprint of the stage's files and inputs; a result made on another stage counts as unknown. The placeholder check judges every mesh a code builder draws, and only the shell builders pass unnamed.

People results in detail (experimental, Section 3.5). The body gates are joins, outline, likeness, triangles and shell. Nev's and Sefa's boot tuck bands lie 9 to 11 mm outside the trouser leg at the bind pose. Scored on drawn surfaces only, no clip of any look is usable: skin seams open and settled cloth crosses itself. Bound as one closed mesh, the player's limbs sink into chest and hips, 87 to 563 crossing points visible from the places' cameras and close views per clip, the armpit surface stretching between arm and chest. The boot tuck bands run through every clip of Nev and Sefa before any motion. Where hair or cloth meets the head or hand skin the two sides are weighted differently and open as the joints move (cloth seams never open), and the settled cloth crosses itself beyond the bind; openings over 10 cm on three looks are not yet confirmed by a render. Binding the take as one closed mesh removes the openings. Pushing points out of the look's own volume, limiting the offset from SOMA-X at the armpits and inner thighs, fold-centred weights and dual-quaternion skinning each left the visible crossings the same or higher. A close-up drawn in a wider A-pose (arms 45 degrees below level, feet 35 cm apart) gives a take with fewer crossings at the bind (the player's 678 against 772) but more in motion, as the clips bring the arms further down: 10,610 visible crossings against 4,078 over the player's 19 clips, and no clip clean in either pose. The clips' joint angles stay within the normal human ranges of the American Academy of Orthopaedic Surgeons, except hip extension of 24 and 27 degrees (limit 20) in two clips.

Decisions recorded in these checks. Three people's gates are decisions of the captain, the agent session that runs world 1, taken on 12 and 13 October 2026 under the creator's rules and open to his reversal. (1) The shell gate counts 99 percent of the soft points rather than every point, with the worst point recorded: the points beyond it were loose cloth edges (cuffs, hems, caps), which no part label names, so they could not be given a joint of their own by label. (2) A look seen only far away is judged on what is drawn far away, not on the near outfit it was made from. (3) The A-pose outline is judged against the finished take, not the raw model, so the take's own loss is not charged to the build. Earlier, the full-body cut (which looks get bodies in full and which only far ones) was set by the people's share of the budget.

C Models and licences

Every model in the framework must allow commercial use of its output (Table C1); we read each licence for its clauses on commercial use. Candidates were dropped under this rule, among them NVIDIA's Lyra 2.0 (weights for internal research only), Hunyuan3D 2.1 (licence does not apply in the European Union) and GarmentCode's fork of NVIDIA Warp (non-commercial research only), which Newton's cloth solver on upstream Warp replaced; Blender's cloth simulation [33] is kept as a second route. The one open doubt is the face LoRA, whose training data is not disclosed.

Table C1. Models in the framework.

Model	Role	Licence
Nano Banana Pro [24]	concepts, close-ups on the prop route and those the open model fails, ground skins	paid API
Nano Banana Pro [24]	finish maps: the close-up drawn again with one flat colour per finish	paid API
Gemini 3.1 Pro [26]	each model's finishes listed on its close-up, and their library surfaces named	paid API
Qwen-Image-Edit-2511 [34]	close-ups, first (before the prop route)	Apache-2.0
Qwen3.8-27B [35]	the close-ups' shape check; each part's material	Apache-2.0
Pixal3D [17]	picture to 3D	MIT; DINOv3 licence [36] for its encoder
PartCrafter [20]	parts	MIT; its non-commercial background remover not used
SegviGen [27]	lays the close-up's finish map onto the 3D shape (surface stage)	MIT, on TRELLIS.2 (MIT); DINOv3 licence for its encoder; its non-commercial background remover and nvdiffrast not used
GeoSAM2 [37]	parts of our own model, seeded by the close-up's regions (trial)	Apache-2.0
Depth Anything V2 Small [38]	ground relief	Apache-2.0
ProcFunc, Infinigen shaders	surfaces	BSD-3-Clause
MOSS-SoundEffect v2.0, CLAP	sound	Apache-2.0
World Labs [39]	optional world step	paid; outputs owned by paid users
FLUX.2 klein 4B [30]	earlier prop pictures; the characters' A-pose pictures from words and their hair's clay pictures	Apache-2.0
MoGe-2 [19], SAM 3 [18]	measuring and finding objects in pictures; counting people and crops on the concepts	MIT; SAM Licence
Kimodo [31]	the characters' motion from sentences	Apache-2.0 code; NVIDIA Open Model License for the Kimodo-SOMA-RP-v1.1 weights, commercial use allowed
LLM2Vec [40] on Meta-Llama-3-8B-Instruct [41]	Kimodo's text encoder, used only while clips are made; nothing from it ships	Llama 3 Community License: commercial use allowed, the 700 million monthly users clause does not apply, "Built with Meta Llama 3" shown, no exclusion in the European Union for the text-only 8B model
SAM 3D Body [42], SAM3DBody-cpp [43]	body shape from a picture	SAM Licence for the weights, SAM3DBody-cpp's converted ONNX and GGUF weights included whatever their tag says; MIT for SAM3DBody-cpp's code
SOMA-X and MHR	the characters' bodies	Apache-2.0
GarmentCode [44], Newton [45] on Warp [46], Blender's cloth [33]	the clothes' sewing patterns and their drape	MIT; Newton and Warp Apache-2.0; Blender is GPL and its output is ours; GarmentCode's Warp fork (non-commercial research only) not used
Hi3DGen [47]	the characters' hair	MIT code; MIT and Apache-2.0 weights

Model	Role	Licence
FLUX.2 klein base 4B [48] with the refcontrol reference-depth LoRA [49]	the characters' faces	Apache-2.0, both; the LoRA's training data is not disclosed
UniRig [50]	the skeleton and skin of a quadruped	MIT code and weights
Robust Skin Weights Transfer [51]	the SOMA-X skeleton's skin weights moved onto a Pixal3D person (a trial)	MIT reference code, ported with trimesh and SciPy; libigl, whose bindings carry GPL-3.0, not used
MakeHuman eyes	the characters' eyes	CC0

The close-ups' shape check was tuned on 40 close-ups that four open picture models drew of the lab's ten objects, each scored by hand, and on Nano Banana Pro's ten pictures of the same objects. With the majority of three sampled answers, the check agreed with the hand score on 29 of the 40, passed all ten of the closed model's pictures and two marginal ones, passed none scored wrong or failed, and failed nine that were good or usable. A single answer at temperature zero was rejected: asked 51 questions twice in one batch, it changed 7 verdicts. On the lab, 7 of the closed model's first pictures drew the room or dimension lines; all 7 passed when redrawn from the crop alone.

Table C2. The close-up stage on world 1. Cloud is rented card time in euros; the closed model's pictures are in dollars.

Run	Close-ups	Open model accepted	Closed model accepted	Left for the creator	Cloud	Closed model
Lab, 22 rows (earlier check: one answer)	22	8	14	0	€4.12	$2.81
World 1's remaining rows (three places)	17	0	10	7	€3.06	$4.02

The picture-to-3D step, the part splitting, the surface bakes, the sound generation, all checks and the assembly run automatically. Coding agents (Claude Code [52]) write the plans, the inventories and every code builder; in the hub's final round, 86 of its 98 models were built in code. Some steps were still done by hand: some generated objects were turned to face the right way by eye, and a few labels were placed by hand before a rule took over that job. The concepts and the verdicts are the creator's own choices.

D The surface library

The surface library holds families of variants, among them painted, steel, aluminium, rubber, plastic, cable, fabric, composite, glass, screen, light, print, wood, ground, foil, water, skin and hair. After the hub's second round it held 71 variants made from 14 ProcFunc recipes; it now holds 369 variants in 26 families made from 22 recipes. Each variant is a recipe with settings, takes its colours only from the palette, and is baked into base-colour, metal-roughness and normal maps at a texture density that depends on how close the player can get. Wear has three levels and comes from causes: edges, and kicks within 32 cm of the floor that a piece stands on. A surface the finish-map route proposes enters an append-only registry and joins the library only when the creator accepts it.

E With and without the agent tools

The tools that check an agent's work were tested on three places of world 1: a room (the workshop, 814 objects), an outdoor place (the street, 1,198 objects) and a dense room (the hangar, 1,657 objects). They are the completion gate (Section 3.6); a triage of the resting check, which judges each object by what holds it up, read from its inventory row, and keeps only real faults [53, 54]; a pick lock, which refuses an export that would show anything but the model the creator picked for a row; and a render-and-compare, which lists the objects least like their close-ups by DINOv2 similarity [16]. In both arms, the same agent (Claude Opus 5.5 in Claude Code, headless) started from the same exported stage and inventory and was asked to make the place as complete as it could in about 50 minutes, through its layout data only; one arm had the framework from before these tools. Each arm ran once, so Table E1 shows what happened, not an average.

Table E1. One session per arm and place. Faults are the resting triage's real faults left on the final stage, measured against the arm's own data; placeholders by the first placeholder check; rows are inventory rows with no object and no written reason. Agent cost is the session's model cost; cloud is its rented machines' cost.

Place, arm	Faults, start → end	Placeholders	Rows missing	Min.	Agent cost	Output tokens	Cloud	Exports	Unknown gate calls
Workshop, without	18 → 13	0 → 0	1 → 0	36	$2.26	30.4k	€0.85	7	not used
Workshop, with	18 → 0	0 → 0	1 → 0	26	$1.42	14.3k	€0.85	2	2 of 4
Street, without	11 → 0	0 → 0	0 → 0	38	$2.40	30.0k	€0.32	9	not used
Street, with	12 → 0	0 → 0	0 → 0	39	$1.21	10.4k	€0.76	1	0 of 5
Hangar, without	14 → 4	5 → 0	0 → 0	47	$2.01	26.6k	€0.00	7	not used
Hangar, with	14 → 9	5 → 0	0 → 0	32	$2.34	29.1k	€0.64	5	3 of 6

The tools made building cheaper more often than they made the scene better. On the workshop and the street, the arm with the tools used about half the agent's output and cost, and exported the stage once or twice instead of seven to nine times, because the triage named the few real faults among hundreds of alarms. On the workshop, 13 of its 18 faults were cleared by writing into the inventory what holds a piece up, which the creator still has to confirm. On the hangar the arm without the tools left fewer faults: the arm with them spent its time replacing five placeholder slabs and trying the drop test, which failed on a room with no ground mesh (fixed after the run), and the nine faults it left are models with stray parts below their feet. No agent claimed a place was complete, so the gate's hook never had to refuse a claim. No row carries an owner's pick, so the pick lock was not tested.

F World 1, place by place

Gap: Figure F1, every place of world 1. For each of the 18 places (the flat, stairwell, street, square and launch; the hub, lab, greenhouse, workshop, airlock, walkway tube, habitat, garage, hangar, wreck and old station; the camp and its grounds), the concept beside the final stage rendered from the concept's own camera. It waits on each place's one assembly passing its completion gate.